

Cortex-M0 Instruction Set Relative & Indexed Addressing Modes

ECE 362

<https://engineering.purdue.edu/ece362/>

More reading

- Textbook, Chapter 6, “Branch and Conditional Execution”, pages 111 – 132.
 - Talking about this in this lecture.
- Textbook, Chapter 5, “Load and Store”, pages 97 – 110.
 - Talking about this in the second half of the lecture.

More reading after that

- Textbook, Chapter 7, "Structured Programming", pages 133 – 160.
 - We'll talk about this soon.
- Textbook, Chapter 8, "Subroutines", pages 161 – 202.
 - We'll talk about this soon.
- ARMv6-M Architecture Reference Manual (436 pages)
 - Get familiar with sections A5.2 (16-bit Thumb encoding) and A6.7 (Alphabetical list of ARMv6-M Thumb instructions)

Instruction Synopsis

Arithmetic

ADDS ADCS SUBS SBCS RSBS
(NEGS) MULS ASRS CMP CMN

Logical

ANDS ORRS BICS EORS MVNS LSLs LSRS RORS TST

Copy Values

MOVS SXTB SXTH UXTB UXTH REV REV16 REVSH

Store

STR STRH STRB STM PUSH

Load

LDR LDRH LDRSH LDRB LDRSB LDM POP ADR

Control Flow

B Bcc BX BL BLX

Exceptions

BKPT WFE WFI SVC NOP

Relative Addressing Mode

- The CPU will fetch/execute instructions one after another. We need a way to tell the CPU to fetch its next instruction from a different address. That means changing the PC to a new value rather than incrementing it.
- We can't easily load a 32-bit value into a register.
- Just as hard to jump. There is no jump instruction with absolute addressing mode like we had with the simple computer or x86.
 - Not enough room in a 16-bit instruction to store a 32-bit address.
 - Instead, we have branches that are relative.

The Relative Branch

- Effectively adding a signed immediate value (+/-) to the PC.
- Unconditional and conditional versions.
- NOTE: During the execution of any instruction, the PC always points to the next instruction.
 - A branch to offset 0 would continue with the next instruction.
 - This is true for every architecture I know of. **Except ARM Cortex**
 - On ARM CPUs, a branch to offset 0 looks like this:



Branch normally used with a label

- We don't have to care about ARM CPUs being strange.
- We don't want to have to hand-compute address offsets... ever.
 - Let the assembler do it for us.

```
.text
.global main
main:    movs    r0, #5           0x2005
        movs    r1, #1           0x2101
L0:      subs    r0, #1           0x3801
        beq     done             0xd000
        b       L2               0xe001
L1:      b       L0              0xe7fb
done:    bkpt    0               0xbe00
L2:      lsls    r1, r1, #1       0x0049
        b       L1              0xe74b
```

“Example 0” on lecture
notes web page

Simpler Branch Example

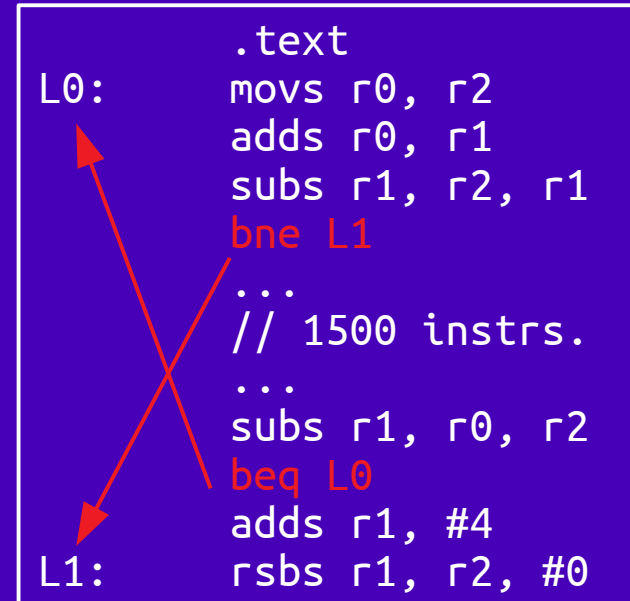
- Now that we can appreciate that the assembler will find branch offsets for us, let's try a simple loop for an example:

```
.text
.global main
main:
    movs r0, #5
    movs r1, #0
loop:
    adds r1, #1
    subs r0, #1
    bne  loop
done:  bkpt
```

“Example 1” on lecture
notes web page

Longer Branch Range Is Better

- There is a limit to how far forward and backward a branch instruction can redirect execution.
- If a program gets too long, it may not be possible to branch to some part of it.
- We usually want to maximize the *range* of a branch instruction.
 - One way to do that is multiply the offset encoded in the instruction by 2.



Branch causes a branch to a target address.

Encoding T1

All versions of the Thumb instruction set.

B<c> <label>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	1	0	1	cond				imm8							

if cond == '1110' then UNDEFINED;

if cond == '1111' then SEE SVC;

imm32 = SignExtend(imm8:'0', 32);

if InITBlock() then UNPREDICTABLE;

Encoding T2

All versions of the Thumb instruction set.

B<c> <label>

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	1	1	0	0	imm11										

imm32 = SignExtend(imm11:'0', 32);

if InITBlock() && !LastInITBlock() then U

Assembler syntax

B{<c>}{<q>} <label>

where:

<c>{<q>} See *Standard assembler syntax fields* on page A6-98.

Note

- Encoding T1 is conditional.
- For encoding T1, <c> must not be AL or omitted.
- For ARMv6-M, for encoding T2, <c> must be AL or omitted.

<label>

The label of the instruction that is to be branched to. The assembler calculates the required value of the offset from the PC value of the B instruction to this label, then selects an encoding that sets imm32 to that offset.

Permitted offsets are even numbers in the range -256 to 254 for encoding T1 and -2048 to 2046 for encoding T2.

Operation

```
if ConditionPassed() then
    EncodingSpecificOperations();
    BranchWritePC(PC + imm32);
```

BNE "PC+12"
1101 0001 00000110
0xD106

B "PC-2"
11100 11111111111
0xE7FF

Instructions always start
on an even address.
So multiply "imm" value
by 2 to get byte offset.

cond	Mnemonic extension	Meaning	Condition flags
0000	EQ	Equal	Z == 1
0001	NE	Not equal	Z == 0
0010	CS ^a	Carry set	C == 1
0011	CC ^b	Carry clear	C == 0
0100	MI	Minus, negative	N == 1
0101	PL	Plus, positive or zero	N == 0
0110	VS	Overflow	V == 1
0111	VC	No overflow	V == 0
1000	HI	Unsigned higher	C == 1 and Z == 0
1001	LS	Unsigned lower or same	C == 0 or Z == 1
1010	GE	Signed greater than or equal	N == V
1011	LT	Signed less than	N != V
1100	GT	Signed greater than	Z == 0 and N == V
1101	LE	Signed less than or equal	Z == 1 or N != V
1110 ^c	None (AL) ^d	Always (unconditional)	Any

a. HS (unsigned higher or same) is a synonym for CS.

b. LO (unsigned lower) is a synonym for CC.

c. This value is never encoded in any ARMv6-M Thumb instruction.

d. AL is an optional mnemonic extension for always.

How to use conditional branches

- In years past, we used to go into tedious detail about how condition codes are set, and how branches interpret those codes. It's simpler than that:
 - Use an instruction that modifies or tests a value.
 - When you want to compare two integers, that's subtraction.
 - An instruction that sets the flags has an 'S' suffix. (e.g. SUBS)
 - Use a branch that interprets the test in the way you want.
 - Branches conditions are set up to work with subtraction.

Branch Selection Example #1

- Subtract R5 from R2 to give R6. ($R6 = R2 - R5$)
 - I want to branch to **error** if $R2 \leq R5$.
 - In other words, if R6 is negative or zero.
 - I want to interpret all values as signed, 2's complement integer representation.

SUBS	R6, R2, R5
BLE	error

1101	LE	Signed less than or equal	$Z == 1$ or $N != V$
------	----	---------------------------	----------------------

Branch Selection Example #2

- Subtract R5 from R2 to give R6. ($R6 = R2 - R5$)
 - I want to branch to **error** if $R2 \leq R5$.
 - In other words, if R6 is negative or zero.
 - I want to interpret all values as **unsigned**, 2's complement integer representation.

SUBS	R6, R2, R5
BLS	error

1001	LS
------	----

Unsigned lower or same

$C == 0$ or $Z == 1$

Branch Selection Example #3

- I don't want to subtract anything.
 - I want to branch to **error** if $R2 > R5$.
 - In other words, if subtraction would be greater than zero.
 - I want to interpret all values as signed, 2's complement integer representation.

CMP	R2, R5
BGT	error

1100	GT	Signed greater than	$Z == 0$ and $N == V$
------	----	---------------------	-----------------------

Wait. Why is it "CMP"
rather than "CMPS"?

Branch Selection Example #4

- I don't want to subtract anything.
 - I want to branch to **error** if $R2 \geq R5$.
 - In other words, if subtraction would be greater or equal to zero.
 - I want to interpret all values as **unsigned**, 2's complement integer representation.

CMP	R2, R5
BHI	error
BEQ	error

1000	HI	Unsigned higher	C == 1 and Z == 0
0000	EQ	Equal	Z == 1

Signed or Unsigned?

- This only matters when you are working with very large integers. For instance: 0xffffffff
 - Is this a negative number?
 - It depends on whether we choose to *interpret* it as a signed or unsigned number.
 - If we want to treat it as a signed integer, it is -1
 - If we want to treat it as unsigned, it is 4,294,967,295

Signed or Unsigned?

- If we do math on numbers that we want to treat as unsigned, it means we need to change the interpretation of “less than zero”.
 - ADDS, SUBS, and CMP work the same either way.
 - The proper branch must be selected to reflect what we want.

Signed or Unsigned?

- In C, we can explicitly declare variables that are never negative (always positive or zero) like this:


```
unsigned int    big32 = 0xffffffff; // AKA uint32_t
unsigned short big16 = 0xffff;      // AKA uint16_t
unsigned char   big8  = 0xff;       // AKA uint8_t
```

- We'll look at this again when we talk about converting C into assembly language.

Logical Shift Left

- We can shift a 32-bit value in a register “left” by any number of bits.
 - This is the same as the C language operator: <<
 - Bits that are shifted out the left end (MSB) of the register are moved into the APSR Carry flag, so we call it “LSLSu”.
 - Zeros are shifted in on the right side (LSB).

Example: LSLS R0, R0, #3



We can use LSL to test bits

- Shift “leftmost” bit into the carry flag.
- How do we test if the carry flag is set?
 - BCS: “Branch if Carry Set”

```
// Count the number of '1' bits in R2.  
// Store result in R0.  
        movs r0, #0  
more:    lsls r2, r2, #1  
        bcs one  
        beq done  
        b more  
one:     adds r0, #1  
        b more  
done:    bkpt
```

0010	CS a	Carry set	C == 1
a. HS (unsigned higher or same) is a synonym for CS. b. LO (unsigned lower) is a synonym for CC.			


“Couldn’t we have used this a few slides ago instead of BHI and BEQ?”

Yes! But you wouldn’t have learned as much.

“Example 2” on lecture notes web page

Other Shift & Rotate Instructions


Logical Shift Right: LSRS R0, R0, #2



Initial R0

New R0


Arithmetic Shift Right: ASRS R0, R0, #2



Initial R0

New R0

Rotate Right: RORS R0, R0, #6



Initial R0

New R0

Bitwise Logical Instructions

ANDS R0, R1

10111001100010011001101000111111

00001000000011111111100000111100 (mask)

000010000001001100110000111100

Initial R0
AND
R1
New R0

ORRS R0, R1

10111001100010011001101000111111

00001000000011111111100000111100 (bits to set)

10111001100011111111010001111111

Initial R0
OR
R1
New R0

BICS R0, R1

10111001100010011001101000111111

00001000000011111111100000111100 (bits to clear)

10110001100000000000010000000011

Initial R0
AND NOT
R1
New R0

Bitwise Logical Instructions

EORS R0, R1

10111001100010011001101000111111

00001000000011111111100000111100 (bits to toggle)

101100011000011001100010000000011

Initial R0
AND
R1

New R0

MVNS R0, R1 // Also known as NOT

10111001100010011001101000111111

01000110011101100110010111000000

Initial R1

New R0

Loads and Stores

Indirect/Indexed Address Mode

- The Cortex-M0 CPU cannot reference memory with arithmetic or logical instructions.
- It can only reference memory with load and store instructions.
 - That's why it's called a load/store architecture.
- It also cannot refer to an absolute address.
 - Instead, it refers to an address *relative* to a register value.
 - A branch is relative to the PC.
 - A load or store is relative to one of registers R0 – R7.

LDR: Load a word of memory into a register

- Let's say R0 holds the value 0x20000000.
- We want to load the four-byte value contained in addresses 0x20000000 – 0x20000003 into register R1.
- This is the instruction:
 - LDR R1, [R0]
- “Use the value in R0 as an address. Read 4-bytes starting from that address, and put the result in R1.”
- **Caveat: R0 must be evenly divisible by 4.**
 - If not, it will invoke a fault handler. You don't want that.

STR: Store register into a word of memory

- Let's say R0 holds the value 0x20000000.
- We want to store the four-byte value contained in R1 to the addresses 0x20000000 – 0x20000003.
- This is the instruction:
 - STR R1, [R0]
- “Use the value in R0 as an address. Write the value of R1 into the 4-bytes starting at that address.”
- **Caveat: R0 must be evenly divisible by 4.**
 - If not, it will invoke a fault handler. You don't want that.

Example of Store / Load

- Consider the following program:

```
ldr r0, =0x200000000 // R0 = 0x200000000
```

```
str r1, [r0]          // mem[0x200000000] = R1
```

```
ldr r1, [r0]          // R1 = mem[0x200000000]
```

Simple Store/Load Example

```
.text
.global main
main:
    movs    r0, #5
    movs    r1, #0
    ldr     r2, =0x20000000

loop:
    str     r0, [r2]
    adds    r1, r0, #1
    ldr     r1, [r2]
    subs    r0, #1
    bne     loop
done:      bkpt
```

- In this example, we initialize R2 to an explicit address (which happens to be the lowest address for SRAM on STM32).

“Example 3” on lecture
notes web page

Another Store/Load Example

```
.text
.global main
main:
    movs r0, #5
    movs r1, #0
    ldr  r2, =storage

loop:
    str  r0, [r2]
    adds r1, r0, #1
    ldr  r1, [r2]
    subs r0, #1
    bne  loop
done:  bkpt

.data
...
storage: .word 0
```

- In this example, we initialize R2 to a label that refers to a location where space is allocated. We don't know where it gets put. The assembler decides for us!

“Example 4” on lecture notes web page

Indirect addressing

- When you see [] around something, it means to treat it as an address and do a memory operation.
 - We don't need to only use R0 for an address. We can use R0 – R7.
 - But we've seen in a past lecture that it's hard work to initialize a register to a 32-bit address.
 - Is there a way we can use LDR/STR to refer to memory *near* that address?

Indexed Addressing Mode

- LDR/STR take an extra *offset* or *index*
- LDR R1, [R0, #12]
- “Take the value of R0, add 12 to it, use the sum as an address, go to that address, read four contiguous bytes, and put it in R1.”
- The offset can be any multiple of 4 from 0 – 124.
- Similar specification for STR.
- **Caveat: R0+offset must be evenly divisible by 4.**

Indexing with a register

- I know what you're going to ask:
“Can I index with a **register** instead of an immediate value?”
- LDR R1, [R0, R2]
“Take the value of R0, add it to the value of R2, use the sum as an address, read four contiguous bytes, and place them in R1.”
- Similar specification for STR.
- **Caveat: R0+R2 must be evenly divisible by 4.**

Also byte/halfword load/store

- There are also
 - STRH: store a halfword (2 bytes) [index is a reg or an #imm multiple of 2 from 0 – 62]
 - LDRH: load an unsigned halfword (2 bytes) [index is the same as STRH]
 - LDRSH: load a signed halfword (sign-extend bit 15) [index must be a reg]
- **Caveat: Halfword load/stores must use an address that is evenly divided by 2.**
 - STRB: store a byte [index is a reg or #imm from 0 – 31]
 - LDRB: load an unsigned byte [index is the same as STRB]
 - LDRSB: load a signed byte (sign-extend bit 7) [index must be a reg]
- A byte load or store has no alignment requirements.

Why are there two kinds of loads?

- Why have LDRH / LDRSH and LDRB / LDRSB when we only have STRH and STRB?

Store	STR	STRH	STRB		
Load	LDR	LDRH	LDRSH	LDRB	LDRSB

- Registers are always treated as 32 bit values.
 - There is no instruction to “add only the lower 16 bits of registers Rx and Ry”
- When loading an 8-bit value from memory, we need to know if it is supposed to be treated as signed or unsigned. (i.e., should bit 7 indicate it is negative)
 - If we load 0xff from memory using LDRB, it will set the destination register to 0x000000ff
 - If we load 0xff from memory using LDRSB, it will set the destination register to 0xffffffff.

Store Byte/Load Byte Example

```
.text
.global main
main:
    movs    r0, #0xfc
    ldr     r2, =0x20000000
    movs    r3, #0

loop:
    strb    r0, [r2]
    // must use reg index
    ldrsb   r0, [r2, r3]
    adds    r0, #1
    bne     loop

done:     bkpt
```

- R0 starts out as 0xfc
- One byte is stored in memory.
- It is loaded from memory as a *signed* byte and then incremented.

“Example 5” on lecture
notes web page

Reading and writing arrays

- How can we implement the following C code:

```
int array1[100];
int array2[100];
...
for(n=0; n<100; n++)
    array1[n] = array2[n] + 5;
```

*For the next few minutes, assume that:
n is just register R0,
R1 is magically the address for array1,
R2 is magically the address for array2.*

```
.data
array1: .space 400
array2: .space 400

.text:
add5:  movs r0, #0           // n = 0
check: cmp r0, #100         // n < 100?
      bge done              // if not, done
      lsls r3, r0, #2       // r3 = r0 * 4
      ldr r4, [r2,r3]       // array2[n]
      adds r4, #5           // add 5
      str r4, [r1,r3]       // array1[n]
      adds r0, #1           // n = n + 1
      b check               // again!

done:  b somewhere_else
```

How to load literal values

- We can jump (update the PC) to a relative offset from the current PC value.
- We can load a value from a relative offset from the current PC value.
- LDR R1, [PC, #12]
“Take the current PC value, round it up to a multiple of 4, add 12 to it, use it as an address, read 4 contiguous bytes from memory, and put the result in R1.”
- Can we store to an address relative to the PC?
 - No. The PC normally points to read-only memory where the instructions are.
- **Bonus: PC (rounded up) +12 is automatically evenly divisible by 4.**
 - But the data you refer to might not be. More on that in a bit.
- We almost always use the preferred form LDR R1, label
- Caveat: The item must be no further than 1020 bytes beyond the PC in the **text** segment.

Example of literal load

```
.text
LDR R1, bigvalue
LDR R2, value2
ADDS R3, R1, R2
...
```

```
.balign 4
bigvalue: .word 0xfedcba98
value2:   .word 0xdecaf123
```



Maximum offset: 1020 bytes.
Word must be in the text segment.
Instructions are 2 bytes long.
Word must be 4-byte aligned.
How do we guarantee **alignment**?

Label is an address

- What value is put into the memory words allocated for `addr1` and `addr2`?

```
.data
0x20000000 array1: .space 400
0x20000190 array2: .space 400
```

array1 is the location in memory of a 400-byte reservation that is uninitialized.

Maybe array1 is in the data segment.

```
.text
0x08000400 addr1: .word array1 → 0x20000000
0x08000404 addr2: .word array2 → 0x20000190
```

addr1 is the location in memory of a 4-byte reservation initialized to be the address of array1.

Values like this in the text segment are called a “literal pool”

Maybe addr1 is in the text segment.

Reading and writing arrays

- Consider this C code again:

```
int array1[100];
int array2[100];
...
for(n=0; n<100; n++)
    array1[n] = array2[n] + 5;
```

```
.data
array1: .space 400
array2: .space 400
```

```
ldr r1, addr1
ldr r2, addr2
add5:  movs r0, #0      // n = 0
check: cmp r0, #100    // n < 100?
      bge done        // if not, done
      lsls r3, r0, #2  // r3 = r0 * 4
      ldr r4, [r2,r3]  // array2[n]
      adds r4, #5      // add 5
      str r4, [r1,r3]  // array1[n]
      adds r0, #1      // n = n + 1
      b check         // again!
done:  b somewhere_else
      .balign 4
addr1: .word array1
addr2: .word array2
```


If that still seems like too much work...

- Just use LDR with a =symbol, and it will automatically build a literal pool at the end of the block of code...

```
.data
array1: .space 400
array2: .space 400
```

```
ldr r1, =array1
ldr r2, =array2
add5:  movs r0, #0          // n = 0
check: cmp r0, #100        // n < 100?
      bge done            // if not, done
      lsls r3, r0, #2     // r3 = r0 * 4
      ldr r4, [r2,r3]     // array2[n]
      adds r4, #5         // add 5
      str r4, [r1,r3]     // array1[n]
      adds r0, #1         // n = n + 1
      b check            // again!
done:  b somewhere_else
```