

# Assembly Language Programming

ECE 362

<https://engineering.purdue.edu/ece362/>

# Reading you should have done

- Textbook, Chapter 7, "Structured Programming", pages 133 – 160.
  - Today's topic.
- Textbook, Chapter 8, "Subroutines", pages 161 – 202.
  - Also today.

# Upcoming reading

- Textbook, Chapter 10, “Mixing C and Assembly”, pages 215 – 236.
  - We’ll talk about this soon.
- Textbook, Chapter 14, “General Purpose I/O (GPIO)”, pages 341 – 372.
  - We’ll talk about this in the next lecture module.
- Textbook, Chapter 11, “Interrupts”, pages 237 – 268.
  - We’ll talk about this later on.

# Assembler Directives

- .cpu cortex-m0
- .fpu softvfp
- .syntax unified
- .thumb
- .data
- .text
- .equ *name*, *replacement*
- .byte n [, n, n, n, ...]
- .hword n [, n, n, n, ...]
- .word n [, n, n, n, ...]
- .space S
- .string "..."
- .align B
  - .balign B
  - .b2align B
- .global symbol

Limit instructions to those recognized by Cortex-M0.

We don't have a floating-point hardware.

Use unified syntax.

Use 16- and 32-bit Thumb instruction set (not 32-bit ARM instructions).

Put following items in read-write memory.

Put following items in read-only memory.

Replace *name* with *replacement*. (like #define)

Don't forget comma.

Reserve and initialize 1 byte of storage for each element in a list of 8-bit integers.

Reserve and initialize 2 bytes of storage for each element in a list of 16-bit integers.

Reserve and initialize 4 bytes of storage for each element in a list of 32-bit integers.

Reserve S bytes of storage. Do not initialize.

Reserve space for the characters plus a null byte terminator.

One of either of the following:

Align the following item on a B-byte boundary.

Align the following item on a 2<sup>B</sup>-byte boundary.

Make the given symbol visible to outside modules.

# Think in C. Code in Assembly.

- We're going to start writing major programs in assembly language.
- The easiest way to do that (that I know of) is to think about what you want in C, and then translate it into assembly language.
- C becomes the modeling language.
- There are three things we need to understand:
  - How to represent variables and arrays
  - How to represent control flow (if, while, do, for)
  - How to represent subroutines

# C-to-Assembly Translation: Variable Update Via Register

```
int x;  
x = x + 1;
```

```
ldr r0, =x          // load addr of x  
ldr r1, [r0]         // load value of x  
adds r1, #1          // add one to value  
str r1, [r0]         // store result to x
```

```
...  
.balign 4            // literal pool  
lit_pool:            // added automatically  
    .word x
```

```
.data  
x: .space 4
```

# C-to-Assembly Translation: Variable Update Via Register

```
int x;  
x = x + 1;
```

```
ldr r0, lit_pool // load addr of x  
ldr r1, [r0]      // load value of x  
adds r1, #1       // add one to value  
str r1, [r0]      // store result to x
```

```
...  
.balign 4          // literal pool  
lit_pool:          // if we added, manually  
    .word x
```

```
.data  
x: .space 4
```

# C-to-Assembly Translation: Statements

```
x = alpha + beta * gamma;    ldr r0, =alpha           // load addr of alpha
                              ldr r1, [r0]           // load value of alpha
                              ldr r0, =beta          // load addr of beta
                              ldr r2, [r0]           // load value of beta
                              ldr r0, =gamma          // load addr of gamma
                              ldr r3, [r0]           // load value of gamma
                              muls r2, r3            // r2 = r2 * r3
                              adds r1, r1, r2         // r1 = r1 + r2
                              ldr r0, =x             // load addr of x
                              str r1, [r0]           // store result to x

.data
.balign 4                    ...
alpha: .space 4              .balign 4 // literal pool added automatically
beta:  .space 4              alpha_addr: .word alpha
gamma: .space 4              beta_addr:  .word beta
x:     .space 4              gamma_addr: .word gamma
                              x_addr:    .word x
```



# C-to-Assembly Translation: Statements

```
x = alpha + beta * gamma;
```

Same as before  
but manual creation  
of the literal pool...

```
.data
.balign 4
alpha:  .space 4
beta:   .space 4
gamma:  .space 4
x:      .space 4
```

```
ldr r0, alpha_addr // load addr of alpha
ldr r1, [r0]        // load value of alpha
ldr r0, beta_addr   // load addr of beta
ldr r2, [r0]        // load value of beta
ldr r0, gamma_addr  // load addr of gamma
ldr r3, [r0]        // load value of gamma
muls r2, r3         // r2 = r2 * r3
adds r1, r1, r2     // r1 = r1 + r2
ldr r0, x_addr      // load addr of x
str r1, [r0]        // store result to x
...
.balign 4
alpha_addr: .word alpha
beta_addr:  .word beta
gamma_addr: .word gamma
x_addr:     .word x
```

# C-to-Assembly Translation: Statements

```
x = alpha + beta * gamma;
```

Same as before  
but manual creation  
of a single-entry literal pool  
when placement of all  
variables is known.

```
.data
.balign 4
alpha:  .space 4
beta:   .space 4
gamma:  .space 4
x:      .space 4
```

```
ldr r0, vars          // load addr of alpha
ldr r1, [r0]           // load value of alpha

ldr r2, [r0, #4]       // load value of beta

ldr r3, [r0, #8]       // load value of gamma
muls r2, r3            // r2 = r2 * r3
adds r1, r1, r2        // r1 = r1 + r2

str r1, [r0, #12]      // store result to x
...
.balign 4
vars: .word alpha
```

# C-to-Assembly Translation: Statements

<code>x = alpha + beta * gamma;</code>	<code>ldr r0, =alpha</code>	<code>// load addr of alpha</code>
	<code>ldr r1, [r0]</code>	<code>// load value of alpha</code>
	<code>ldr r2, [r0, #4]</code>	<code>// load value of beta</code>
	<code>ldr r3, [r0, #8]</code>	<code>// load value of gamma</code>
	<code>muls r2, r3</code>	<code>// r2 = r2 * r3</code>
	<code>adds r1, r1, r2</code>	<code>// r1 = r1 + r2</code>
	<code>str r1, [r0, #12]</code>	<code>// store result to x</code>

Same as before  
but automatic creation  
of the literal pool  
when placement of all  
variables is known.

```
.data
.balign 4
alpha: .space 4
beta:  .space 4
gamma: .space 4
x:     .space 4
```

# C-to-Assembly Translation: "if-then-else"

```
if (expr) {  
    then_statements; ...  
} else {  
    else_statements; ...  
}
```

```
if1:      expr  
          ...  
          branch_if_not else1  
then1:    then_statements  
          ...  
          b endif1  
else1:    else_statements  
          ...  
endif1:
```

# C-to-Assembly Translation: "if-then-else" example

```
if (x > 100) {  
    x = x - 1;  
} else {  
    x = x + 1;  
}
```

```
if1:  
    ldr r0, =x  
    ldr r1, [r0]  
    cmp r1, #100  
    ble else1
```

```
then1:  
    ldr r0, =x  
    ldr r1, [r0]  
    subs r1, #1  
    str r1, [r0]  
    b endif1
```

```
else1:  
    ldr r0, =x  
    ldr r1, [r0]  
    adds r1, #1  
    str r1, [r0]
```

```
endif1:
```

Branch is  
always based  
on the *opposite*  
of the condition  
tested for.

# C-to-Assembly Translation: "do-while"

```
do {  
    do_body_stmts; ...  
} while (expr);
```

```
do1:  
    do_body_stmts  
    ...  
while1:  
    expr  
    branch_if_yes do1  
enddo1:
```

# C-to-Assembly Translation: "do-while" example

```
do {  
    x = x >> 1;  
} while (x > 2);
```

```
do1:  
    ldr r0, =x  
    ldr r1, [r0]  
    asrs r1, r1, #1  
    str r1, [r0]  
while1:  
    ldr r0, =x  
    ldr r1, [r0]  
    cmp r1, #2  
    bgt do1  
enddo1:
```

# C-to-Assembly Translation: "while"

```
while (expr) {  
    while_body_stmts; ...  
}
```

```
while1:  
    expr  
    branch_if_not endwhile1  
do1:  
    while_body_stmts; ...  
    b while1  
endwhile1:
```



# C-to-Assembly Translation: "while" example

```
while (x > y) {  
    y = y + 1;  
}
```

```
while1:  
    ldr r0, =x  
    ldr r1, [r0]  
    ldr r0, =y  
    ldr r2, [r0]  
    cmp r1, r2  
    ble endwhile1  
  
do1:  
    ldr r0, =y  
    ldr r1, [r0]  
    adds r1, #1  
    str r1, [r0]  
    b while1  
  
endwhile1:
```

# C-to-Assembly Translation: "for"

```
for (init; check; next_stmt) {  
    for_body_stmts; ...  
}
```



```
init;  
while (check) {  
    for_body_stmts; ...  
    next_stmt;  
}
```

```
for1:  
    init  
forcond1:  
    check  
    branch_if_not fordone1  
forbody1:  
    for_body_stmts; ...  
fornext1:  
    next_stmt  
    b forcond1  
fordone1:
```

# C-to-Assembly Translation: "for" example

```
for (q=0; n>=d; q++) {  
    n = n - d;  
}  
r = n;
```

```
for1:                                fordone1:
    ldr r0, =q                        ldr r0, =n
    movs r1, #0                       ldr r1, [r0]
    str r1, [r0]                       ldr r0, =r
                                        str r1, [r0]
forcond1:                             ldr r0, =n
    ldr r0, =n                        ldr r1, [r0]
    ldr r1, [r0]                       ldr r0, =d
    ldr r0, =d                        ldr r2, [r0]
    ldr r2, [r0]                       cmp r1, r2
    cmp r1, r2                         blt fordone1
    blt fordone1
forbody1:                             ldr r0, =n
    ldr r0, =n                        ldr r1, [r0]
    ldr r1, [r0]                       ldr r0, =d
    ldr r0, =d                        ldr r2, [r0]
    ldr r2, [r0]                       subs r1, r1, r2
    subs r1, r1, r2                    ldr r0, =n
    ldr r0, =n                        str r1, [r0]
    str r1, [r0]
fornext1:                             ldr r0, =q;
    ldr r0, =q;                       ldr r1, [r0]
    ldr r1, [r0]                       adds r1, #1
    adds r1, #1                       str r1, [r0]
    str r1, [r0]                       b forcond1
    b forcond1
```

# Simple Division

- The ARM Cortex-M0 has no divide instruction.
  - You should feel lucky that it has a multiply.
- Simple way to divide one number by another:
  - Initialize the quotient to be zero.
  - While numerator  $\geq$  denominator,
    - Subtract the divisor (denom.) from the numerator
    - Increment the quotient
  - The dividend (numerator) is now the remainder.


```
for (q=0; n>=d; q++) {  
    n = n - d;  
}  
r = n;
```

# Division Example:

- 90 / 12: quotient = 0
- 78 / 12: quotient = 1
- 66 / 12: quotient = 2
- 54 / 12: quotient = 3
- 42 / 12: quotient = 4
- 30 / 12: quotient = 5
- 18 / 12: quotient = 6
- 6 / 12: quotient = 7
- Done. Quotient = 7, Remainder = 6

"Can't we do this any faster?"

Yes.



# Consider a simple program:

- Let's use subroutines this time.
  - How do we even do this?

```
int x;  
int main(void) {  
    x = x + 1;  
    empty_subroutine();  
    x = x - 2;  
    ...  
}  
  
void empty_subroutine(void) {  
}
```

# The Stack: Subroutines

- The Cortex-M0 has a stack pointer (SP), but no JSR instruction.
  - Recall that JSR on the simple computer:
    - Decrement (subtracted one from) SP
    - Wrote PC of next instruction to [SP]
    - Jumped to immediate argument.
  - Cannot have all that complexity in a RISC CPU.
  - To invoke a subroutine with a Cortex-M0 CPU, we:
    - save pointer to next instruction in Link Register (LR)
    - branch to the destination address
    - Instruction is BL (Branch with Link)
    - It is the callee's responsibility to save the LR

BL: (branch with link)

“Put the current PC  
into the LR register  
and branch to the  
given label.”

# Example of BL

```
ldr    r0,=x
ldr    r1,[r0]
adds   r1,r1,#1
str    r1,[r0]
bl     empty_subroutine
ldr    r0,=x
ldr    r1,[r0]
subs   r1,r1,#2
str    r1,[r0]
bkpt
...
```

```
empty_subroutine:
bx lr
```

```
.data
x: .word 1
```

"Wait, you didn't even use the stack there."

True. This is the smallest example possible.

Example 0 on  
lecture notes page

BX: Branch and Exchange

"Put the contents of the  
specified register into the PC."



# Saving LR on the Stack

- Cortex-M0 has PUSH and POP instructions:

- PUSH:

- Decrement SP (multiple times).
    - Write multiple registers into memory pointed to by SP.
      - One of those registers can be LR.

See ARMv6-M Architecture  
Reference Manual  
Section A6.7.50

- POP:

- Read multiple registers from memory pointed to by SP.
      - One of those registers can be PC.
    - Increment SP (multiple times).

See ARMv6-M Architecture  
Reference Manual  
Section A6.7.49

# A Real Subroutine

```
.syntax unified
.cpu cortex-m0
.thumb

.global main
main:
    movs    r0, #25           // initialize r0 with 25
    bl      incr0             // "call" incr0.  LR is next instruction.
    bkpt    0                 // What is value of r0 when we get here?

incr0:
    push    {lr}              // Save the LR "on the stack"
    adds    r0, #1            // Add 1 to r0
    pop     {pc}              // Pop one word "from the stack" to PC.
```

# More Complex Functions

- Consider this C program:

- main calls first

- first calls second
    - first calls second again

- main calls first again

- first calls second
    - first calls second again

- first and second must know where to return to!

By the way, what's  
the final value of x?

```
int x = 0;
```

```
int main() {  
    first();  
    first();  
    ...  
}
```

```
void first() {  
    second();  
    second();  
}
```

```
void second() {  
    x += 1;  
}
```

# Don't always need to save LR

```
.data
x: .word 0
.text
.global main
main:
```

Example 1 on  
lecture notes page

```
    bl    first
    bl    first
    bkpt
first:
    push  {lr}
    bl    second
    bl    second
    pop   {pc}
```

```
second:
    ldr   r0, =x
    ldr   r1, [r0]
    adds r1, #1
    str   r1, [r0]
    bx    lr
```

first() must save LR  
before calling second()

first() returns with pop {pc}

second() is a “leaf function”.  
It calls no other functions, so  
LR will not be overwritten.  
No need to push it.

```
int x = 0;
```

```
int main() {
    first();
    first();
    ...
}
```

```
void first() {
    second();
    second();
}
```

```
void second() {
    x += 1;
}
```

# We can save more registers with PUSH/POP than just LR/PC

```
.syntax unified
.cpu cortex-m0
.thumb
```

Example 2 (or something similar) on lecture notes page

```
.global main
main:
```

```
    movs    r0, #25           // initialize r0 with 25
    bl      incr0             // "call" incr0.  LR is next instruction.
    nop
    bkpt     0                // What is value of r0 when we get here?
```

```
incr0:
```


```
    push    {r1-r4,r6,lr}     // Save lots of registers to the stack
    adds    r0, #1            // Add 1 to r0
    pop     {r1-r4,r6,pc}     // Pop several words from stack
```

r1-r4 means R1 **through** R4 (R1,R2,R3,R4)

# Be Careful With Stack Memory

- Your development board has 32K (32768) bytes of SRAM (static random-access memory).
  - 0x20000000 - 0x20007fff
- SP is initialized to 0x20008000.
  - (It is decremented by 4 before writing a word to memory.)
- Things like variables in the data segment start at 0x20000000.
- It's really all the same memory.
- If you keep pushing things onto the stack, you will eventually overwrite variables.

0x20008000



0x20000000

# Consider a recursive subroutine:

```
int total = 0;

int sum(int x) {
    if (x > 0) {
        total += x;
        sum(x-1);
    }
}

int main() {
    sum(3);
}
```

```
sum:    push {lr}           // save link register
if1:    cmp  r0,#0          // x > 0
        ble  endif1        // no? goto endif1
        ldr  r1,=total      // get addr of total
        ldr  r2,[r1]        // load from total
        adds r2,r0          // total = total + x
        str  r2,[r1]        // store total
        subs r0,#1          // x - 1
        bl   sum            // sum(x - 1)
endif1: nop
        pop  {pc}           // return

.global main
main:
        movs r0, #3
        bl   sum
        wfi

.data
total   .word 0
```

Example 3 on  
lecture notes page

# Application Binary Interface

- Every platform defines an ABI for functions and subroutines.
- Ours looks like this:
  - The caller passes the first four parameters in R0-R3.
    - Any more parameters go on the stack.
  - The callee is allowed to modify R0-R3 and R12.
    - If any other registers are changed, they must first be saved to the stack.
  - The return value from a function is in R0.
- These are the rules if you want your assembly language functions to get along with C functions.
  - We can do anything we want if we're writing our own code in assembly language, but a C compiler maintains certain expectations. When we link in C code, we have to play by the rules.



# Why have an ABI?

- Clearly defines who needs to save registers and when.
- As long as everyone adheres to the ABI, you know that `fn()` will not mess with R4 – R11.
- `calc()` saves and restores R4 so that any caller can trust it not to modify R4.

```
int calc(int x) {  
    return x + fn(3);  
}
```

```
calc:  
    push {r4,lr}  
    movs r4,r0  
    movs r0, #3  
    bl fn  
    adds r0,r4  
    pop {r4,pc}
```

# Consequences of the ABI

- You almost never want to save and restore R0.
  - Don't push or pop that.
- You almost never need to save and restore R1 – R3 or R12.
- Generally, if you need to use R4 – R11, save them at the beginning of your subroutine and restore them at subroutine exit.

# Check the “Bonus Lectures”

- There are many nuances of subroutines that are difficult for students to understand.
- The course textbook has good coverage of these topics.
- Otherwise, go through the bonus lectures on subroutines and recursive flow control.